

## INDEX

### Parameter sniffing (*cont.*)

- definition, [511](#)
- local variable, [512](#)
  - execution plan, [514–515](#)
  - query, [513–514](#)
- maintenance reexamination, [528](#)
- stored procedure, [515](#)
- sys.dm\_exec\_query\_stats output, [516](#)
- values, [515](#)

### Parse tree, [454](#)

### Partition elimination, [67](#)

### Performance Monitor counters, [693](#)

### Performance tuning process

- baseline performance, [11–13](#)
- data access layer, [13](#)
- database connection, [5](#)
- database design, [5](#)
- hardware and software factors, [3](#)
- high level database, [15](#)
- iteration process
  - costliest query, [8–9](#)
  - user activity, [7](#)
- low level database, [15](#)
- optimization, [6](#)
- performance killers
  - cursors, [21](#)
  - excessive blocking and deadlocks, [18](#)
  - excessive fragmentation, [21](#)
  - frequent recompilation, [20](#)
  - inaccurate statistics, [17](#)
  - inappropriate database design, [19](#)
  - insufficient indexing, [16](#)
  - nonreusable execution plans, [20](#)
  - non-set-based operations, [19](#)
  - parameter sniffing, [18](#)
  - query design, [17](#)
  - SQL Server, [15–16](#)
- vs.* price, [10](#)

### query optimization, [1–2](#), [6](#)

### root causes, [13–14](#)

### SQL Server configuration, [5](#)

### Physical Design Structure (PDS), [282](#)

### Plan cache, [468](#)

### Plan forcing, Query Store, [313–315](#)

### Plan guides

- execution plan, [562](#), [566](#)
- Index Seek operation, [566](#)
- OPTIMIZE FOR query hint, [561–563](#)
- SELECT operator property, [563](#)
- sp\_create\_plan\_guide\_from\_  
handle, [567](#)
- T-SQL statement, [565](#)

### Platform as a Service (PaaS), [799](#)

### Prepared workloads, plan reusability

- prepare/execute model, [499](#)
- sp\_executesql
  - additional output, [498](#)
  - parameterized plan, [496](#)
  - plan sensitivity, [497](#)
  - SELECT statement, [495–496](#)
- stored procedures
  - extended events, [489–490](#)
  - output, [490](#)
  - reuse, [491](#)

### Prepare/execute model, [475](#), [499](#), [506](#)

### Pseudoclustered index, [234](#)

## Q

### Query analysis

- missing statistics issue
  - ALTER DATABASE command, [392](#)
  - CREATE STATISTICS statement, [394](#)
  - execution plan, [395–396](#)
  - graphical plan, [394](#)
  - Index Scan operator, [394](#)

- SELECT statement, 393
- test table creation, 392
- outdated statistics issue
  - Analyze Actual Execution Plan, 399
  - database, 400
  - DBCC SHOW\_STATISTICS
    - command, 397
  - estimated *vs.* actual rows value, 398, 400
  - execution plan, 398
  - FULLSCAN, 399
  - iFirstIndex, 397
  - inaccurate\_cardinality\_
    - estimate, 396
  - SELECT statement, 397
  - Table Scan operator, 398
- Query compilation, 454
- Query design
  - advantage of query store, 909
  - aggregate functions, MIN and MAX, 614-615
  - arithmetic expressions, 903-904
  - compile stored procedure, 909
  - COUNT(\*) and EXISTS, 611-612
  - database cursors, 908
  - database object owner, 902
  - database transactions, 907
  - implicit conversion, 905
  - implicit data type conversion, 608-611
  - local variables in batch query
    - estimated number of rows, 618-619
    - execution plan, 617-618
    - index seek and Key Lookup
      - operators, 617-618
    - I/O output, 619
    - parameter sniffing, 620
    - SELECT statement, 617
    - TransactionHistory table, 618
    - WHERE clause, 616
  - naming stored procedures, 621-623
  - nesting views, 905
  - network round-trips
    - execute multiple queries, 624
    - SET NOCOUNT, 624
  - ORDER BY clause, 616
  - nonsargable search conditions, 902-903
  - optimizer hints, 904
  - reusing execution plans, 906
  - SET NOCOUNT ON command, 902
  - techniques, 607
  - transaction cost
    - atomic action, 624
    - lock overhead, 627-629
    - logging overhead, 625-627
  - transaction log, 905-906
  - UNION ALL, 613-614
  - UNION clause, 612-613- Query design recommendations
  - avoiding optimizer hints
    - (*see* Optimizer hints)
  - domain and referential integrity
    - DRI, 600
    - NOT NULL constraint, 597
  - effective indexes
    - avoid arithmetic operators, 580
    - avoid nonsargable search conditions (*see* Nonsargable search conditions)
    - avoid WHERE clause column (*see* WHERE clause columns)
  - performance, 570
  - small result sets
    - limited number of columns, 571-572
    - WHERE clause, 572-573

## INDEX

- Query execution, [105–106](#)
- Query hash, [468, 499–503](#)
- Query optimization
  - multiple phases (*see* Multiple optimization phases)
  - parallel plan optimization (*see* Parallel plan optimization)
  - simplification, [459](#)
  - statistics (*see* Statistics)
  - steps, [457–458](#)
  - trivial plan match, [459](#)
- Query performance metrics
  - Extended Events (*see* Extended Events sessions)
  - methods, [128](#)
  - sys.dm\_exec\_query\_stats DMO, [129–130](#)
- Query plan, [468, 499–503](#)
- Query processor tree, [454](#)
- Query recompilation
  - compile process, [535](#)
  - deferred object resolution (*see* Deferred object resolution)
  - execution plan, [530, 545](#)
  - Extended Events, [533–535](#)
  - implementation
    - DDL/DML statements, [551](#)
    - disabling automatic statistics update, [556](#)
    - KEEPFIXED PLAN, [554](#)
    - OPTIMIZE FOR query hint, [559, 561](#)
    - plan guides (*see* Plan guides)
    - SET options, [558](#)
    - statistics change, [553](#)
    - table variables, [556](#)
  - index IX\_Test, [531](#)
  - nonbeneficial recompilation, [532](#)
  - RECOMPILE clause (*see* RECOMPILE clause)
  - schema/binding changes, [536](#)
  - SELECT statement, [530, 532](#)
  - SET options, [544](#)
  - sp\_recompile, [545](#)
  - SQL Server rules, [531](#)
  - sql\_statement\_recompile event, [532](#)
  - statement recompilation, [535–536](#)
  - statistics changes, [536](#)
  - stored procedure, [530](#)
- Query Store
  - behavior, [294–296](#)
  - controlling, [306–308](#)
  - monitoring, query performance, [293](#)
  - plan forcing, [313–315](#)
  - query information
    - basic structure, [297–298](#)
    - parameter definition, [299](#)
    - primary file group, [298](#)
    - query\_hash value, [297](#)
    - simple parameterization, [300–301](#)
    - stored procedure, execution plan, [298](#)
    - T-SQL, [299–300](#)
  - reporting
    - AdventureWorks2017 database, [309](#)
    - details of information, [312–313](#)
    - forcing and unforcing plans, [313](#)
    - performance behaviors and execution plans, [312](#)
    - Top 25 Resource Consumers report, [310–311](#)
    - T-SQL, [309](#)
  - runtime data and wait statistics, [301–305](#)
  - safety and reporting mechanism, [293](#)
  - system views, [297](#)
  - upgrades, [316–317](#)

## R

Read Committed isolation level, [661–663](#)  
 Read-only concurrency model, [725–726](#)  
     cost benefits, [734](#)  
     drawback, [735](#)  
 Read Uncommitted isolation level, [660–661](#)  
 Recompilation threshold (RT), [536](#)  
 RECOMPILE clause  
     with CREATE PROCEDURE  
         statement, [547](#)  
     EXECUTE statement, [548–549](#)  
     query hint, [549–550](#)  
 Remote procedure call (RPC), [108](#)  
 Referential integrity, [887](#)  
 Repeatable Read isolation level, [663](#),  
     [665–667](#)  
 Row-level compression, [253](#)  
 Row-level lock, [642–643](#)  
 Rowstore indexes, [257](#)

## S

Sargable search conditions, [574](#)  
 Scalar functions, [586–589](#)  
 Scroll locks concurrency model, [727](#)  
     benefit, [736](#)  
     cost overhead, [736](#)  
 Semantic search, [267](#)  
 Serializable isolation level, [679–680](#)  
     bonus payment, [668](#)  
     business functionality, [668](#)  
     clustered index, [673](#)  
     HOLDLOCK locking, [670](#)  
     PayBonus transaction, [669](#)  
     phantom read, [667](#)  
     SELECT statement, [667](#)  
     side effects, [672](#)  
     sys.dm\_tran\_locks, [671–673](#)

Server-side cursors  
     characteristics, [725](#)  
     cost benefits, [733](#)  
     cost overhead/disadvantages, [734](#)  
 SET statement, [544–545](#)  
 Simple parameterization  
     autoparameterized plan, [482–483](#), [485](#)  
     definition, [481](#)  
     limits, [485](#)  
     using template, [484](#)  
 Snapshot isolation, [674](#)  
 Solid-state disks (SSDs), [51](#)  
 Solid-state drives (SSD), [36](#), [62](#)  
 Spatial index, [267](#)  
 sp\_executesql technique  
     additional output, [498](#)  
     definition, [475](#)  
     and EXECUTE, [507–508](#)  
     parameterized plan, [496](#)  
     plan sensitivity, [497](#)  
     SELECT statement, [495–496](#)  
 SQL query performance  
     actual *vs.* estimated execution plans, [161](#)  
     costly query  
         multiple executions, [134](#), [137](#)  
         query execution plan, [133](#)  
         query optimizer, [133](#)  
         reads fields, [133–134](#)  
         single execution, [134–135](#)  
         slow-running query, [141–142](#)  
     execution plan  
         Active Expensive Queries, [173](#)  
         Activity Monitor, [173](#)  
         actual *vs.* estimated execution  
             plans, [142–143](#)  
         Clustered Index Scan operator, [145](#)  
         compare plans, [166](#)  
         cost-effective execution plan, [142](#)

SQL query performance (*cont.*)

- dynamic management views and functions, [164](#)
- Find Node, [165](#)
- graphical execution plan, [142–143](#)
- identification, [148, 150](#)
- index effectiveness, [150](#)
- index scan/seek, [145–146](#)
- Live Query Statistics, [170](#)
- operator selection properties, [147–148](#)
- physical and logical operation, [147](#)
- query optimizer, [142](#)
- query window, [171](#)
- scenarios, [169](#)
- SET SHOWPLAN\_XML
  - command, [143](#)
- Show Execution Plan XML, [144](#)
- SQL Server Management Studio
  - 2016, [172](#)
  - tooltip sheet, [146–147](#)
  - XML execution plan, [142–143](#)
- Extended Events, [132](#)
- join effectiveness
  - adaptive, [159](#)
  - hash, [154](#)
  - merge, [156](#)
  - nested loop, [157](#)
- parameters, [133](#)
- plan cache, [164](#)
- query resource cost
  - client statistics, [176](#)
  - execution time, [178](#)
  - Extended Events, [176](#)
  - QueryTimeStats and Wait Stats, [182](#)
  - STATISTICS IO, [179](#)
- query thread profiles, [174](#)

- reducing database blocking and pressure, [132](#)

SQL Server Management Studio 2016, [172](#)

SQL server optimization

- configuration settings
  - ad hoc workloads, [895](#)
  - blocked process threshold, [895](#)
  - cost threshold, [894](#)
  - database file layout, [895–896](#)
  - data compression, [896](#)
  - max degree of parallelism, [894](#)
  - memory configuration, [894](#)
- database administration
  - AUTO\_CLOSE, [898](#)
  - AUTO\_SHRINK, [899](#)
  - minimum index
    - defragmentation, [898](#)
  - up-to-date statistics, [897](#)
- database backup
  - compression, [901](#)
  - distribution, [900](#)
  - transaction log frequency, [899](#)
- database design (*see* Database design)
- query design (*see* Query design)
- Statement recompilation, [535–536](#)
- Static cursors
  - characteristics, [729](#)
  - cost benefits, [738](#)
  - cost overhead, [738](#)
- Statistics
  - auto create statistics, [401](#)
  - auto update statistics, [401](#)
  - backward compatibility, [400](#)
  - cardinality estimation
    - AND calculation, [370, 372](#)
    - enabling and disabling, [377](#)
    - Estimated Row Count value, [371](#)
    - Find Node menu selection, [376](#)

- FULLSCAN, 373
- graphical execution plan, 376
- query\_optimizer\_estimate\_
  - cardinality event, 373–374
- OR calculation, 370, 372
- stats\_collection\_id values, 375
- data retrieval strategy, 360–361
- definition, 337
- density, 363
- DMOs, 379
- filtered index, 367
- histogram, 359–360, 362
- iFirstIndex, 361
- maintenance behavior
  - auto create statistics, 381–382, 391
  - automatic maintenance, 381
  - auto update statistics, 382, 385
  - management setting, 387
  - manual maintenance, 385
- maintenance status, 390
- multicolumn index, 364
- nonindexed column
  - ad hoc T-SQL activities, 348
  - ALTER DATABASE command, 348
  - AUTO\_CREATE\_STATISTICS
    - process, 352
  - cost comparison, 357
  - execution plan, 353
  - Index Seek operation, 354
  - query optimizer, 347
  - SQL Server, 348, 355
  - sys.stats table, 352
- nonindexed columns
  - AUTO\_CREATE\_STATISTICS
    - OFF, 355–357
  - AUTO\_CREATE\_STATISTICS
    - ON, 351
  - auto\_stats events, 351
- DATABASEPROPERTYEX
  - function, 350
- data distribution, 348
- FROM clause, 356–357
- graphical plan, 358
- missing\_column\_statistics, 359
- query cost, 357
- sample tables, 350
- SELECT statement, 350
- sys.stats\_columns system, 354
- Test1.Test1\_C2 and Test2.Test2\_C2, 352–353
- query analysis (*see* Query analysis)
- query optimization
  - auto\_stats events, 345
  - Extended Events, 341
  - indexed column, 338–339
  - large data modification, 343–344
  - nonclustered index, 340
  - outdated statistics, 345
  - SELECT statement, 340, 343
  - small data modification, 341–342
  - system resources, 339
- sampling rate, 402
- small and large result set
  - queries, 362
- Storage area network (SAN), 53, 61
- Stored procedures
  - definition, 475
  - extended events, 489–490
  - first execution, 492
  - output, 490
  - performance benefits, 492–494
  - reuse, 491
  - UserOne user, 494
- Syntax-based optimization, 455–456
- system\_health Extended Events session, 121–123

## INDEX

### T

Table-level lock, [646](#)

T-SQL cursors, [746–750](#)

### U

UNIQUE constraint, [886](#)

### V

Virtual machine (VM), [24](#), [87–88](#)

### W, X, Y, Z

WHERE clause

columns, [572–573](#)

custom scalar UDF,

[586–589](#)

date part comparison

CONVERT function, [585](#)

DATEPART function, [584](#)

DATETIME column, [583](#)

SUBSTRING *vs.* LIKE, [582](#)